


Compilers

Topic: The View from 35,000 Feet
Monsoon 2011, IIIT-H, Suresh Purini

ACK: Some slides are based on Keith Cooper's CS412 at Rice University

High Level View of a Compiler

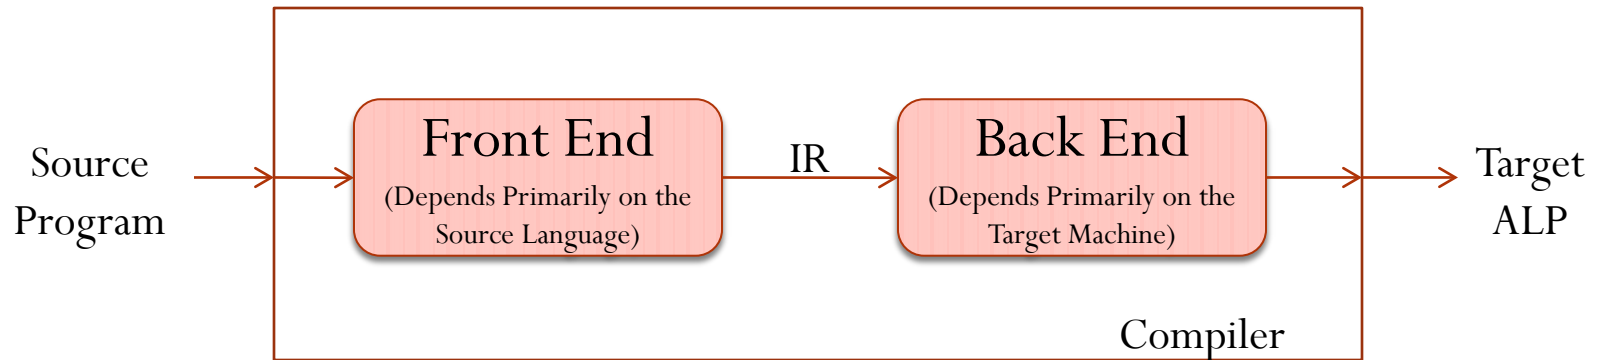


Implications

- Must recognize legal (and illegal) programs
- Must generate correct and efficient code
- Must manage storage of all variables (and code)
- Must agree with OS & linker on format for object code

Big step up from assembly language—use higher level notations

Traditional Two-pass Compiler



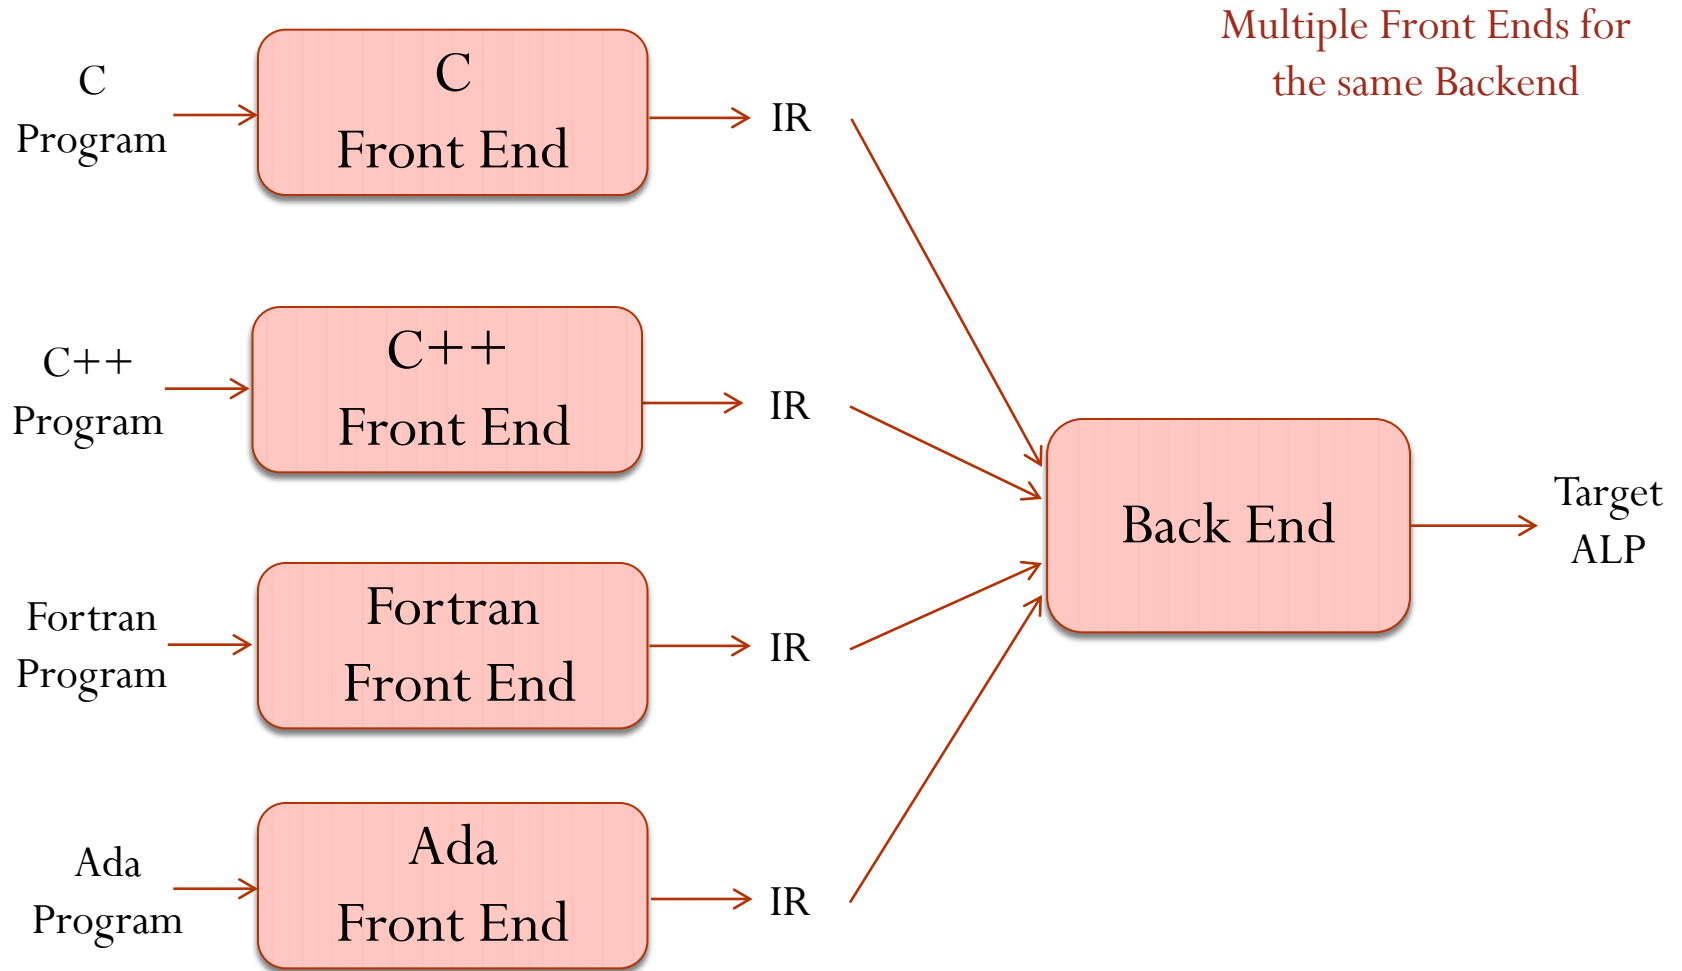
Implications

- Use an intermediate representation (IR)
- Front end maps legal source code into IR
- Back end maps IR into target machine code
- Admits multiple front ends & multiple passes


Classic principle from software engineering:
Separation of concerns

Typically, front end is $O(n)$ or $O(n \log n)$, while back end is NPC


Why IR?



Why IR?



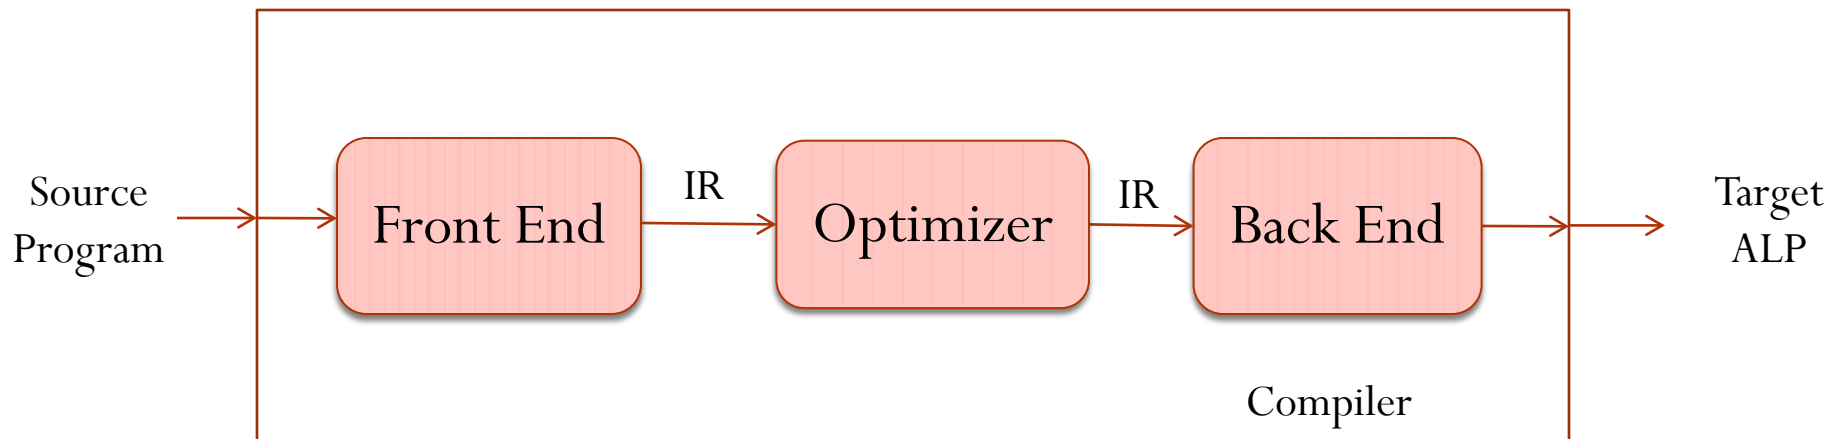
A Common Fallacy



Can we build $n \times m$ compilers with $n+m$ components?

- Must encode all language specific knowledge in each front end
- Must encode all features in a single IR
- Must encode all target specific knowledge in each back end
- Successful in systems with assembly level (or lower) IRs
 - e.g. gcc or llvm

Structure of a Three Phase Compiler



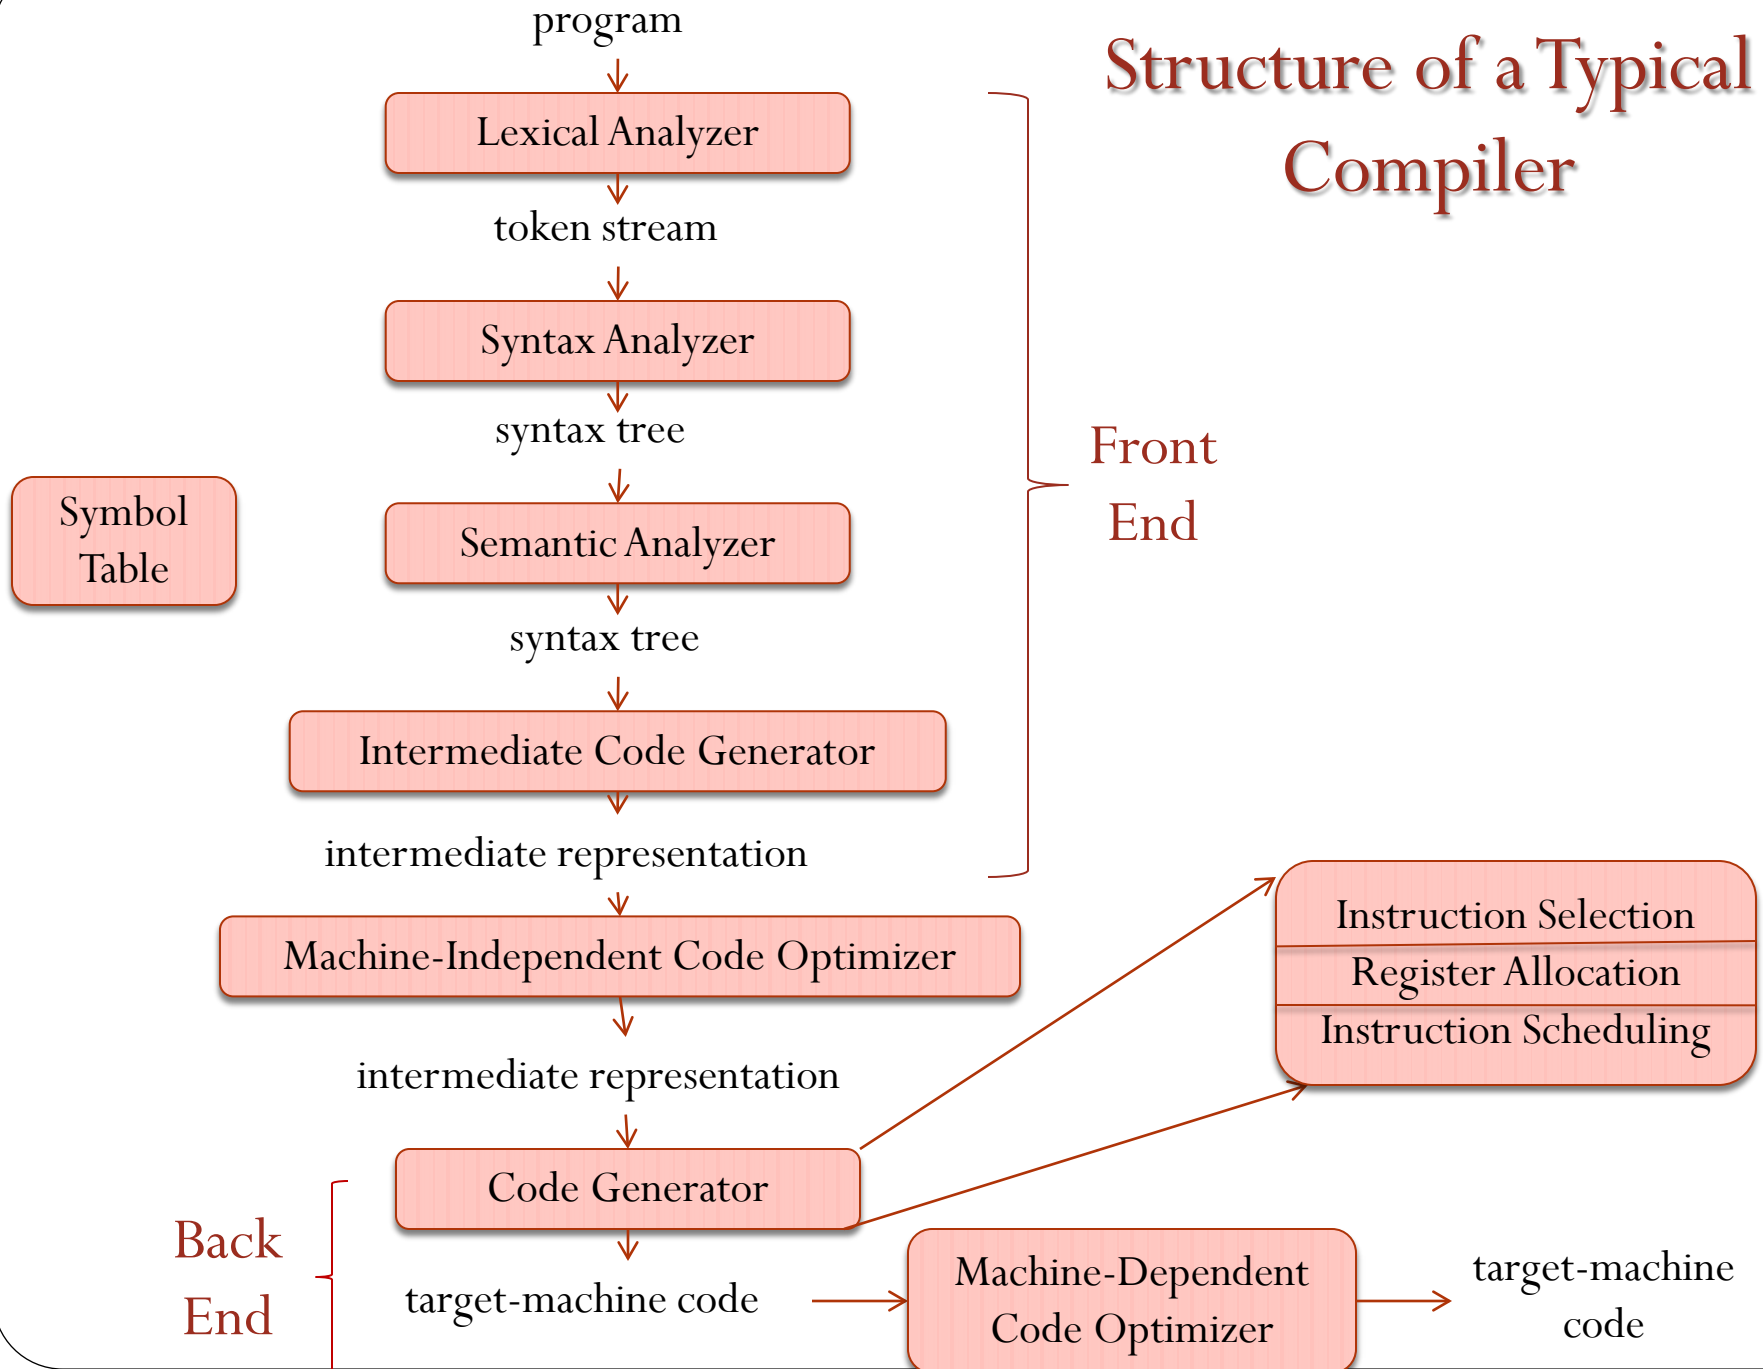
An Optimizer

- Analyzes IR and rewrites (or transforms) IR
- Primary goal is to reduce running time of the compiled code
 - May also improve space, power consumption, ...
- Must preserve “meaning” of the code
 - Measured by values of named variables

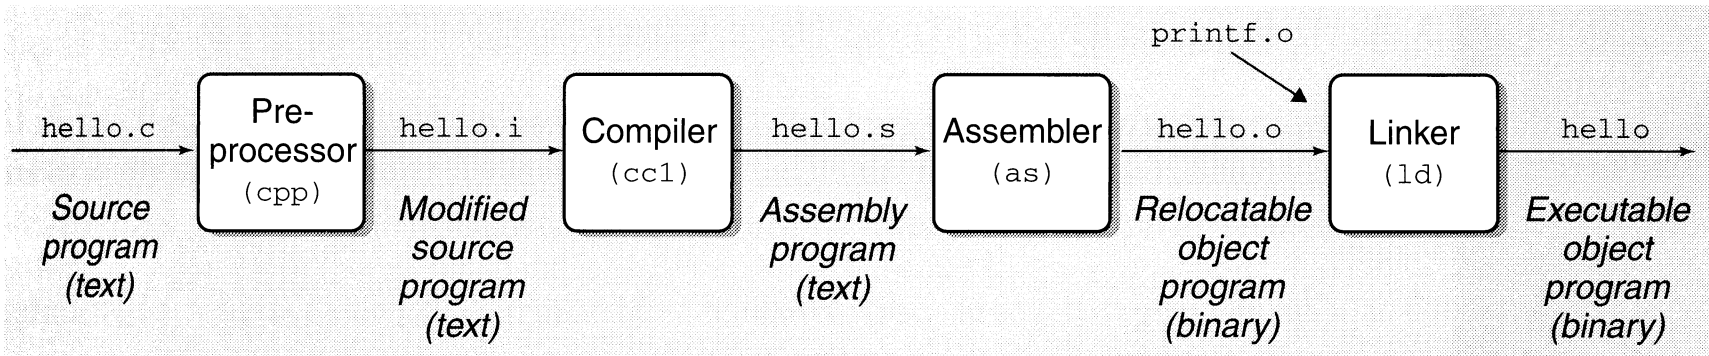
Levels of IRs

- Multiple optimization passes can be made over the IR improving the final program performance.
- Typically multiple IRs at various phases of compilation
 - High level IRs
 - Abstract Syntax Trees – Useful for loop transformations, procedure inlining etc.
 - Intermediate level IRs
 - Three address code, Static Single Assignment form – Useful for optimizations like constant propagation etc.
 - Low level IRs – Good for Low level machine dependent code optimizations

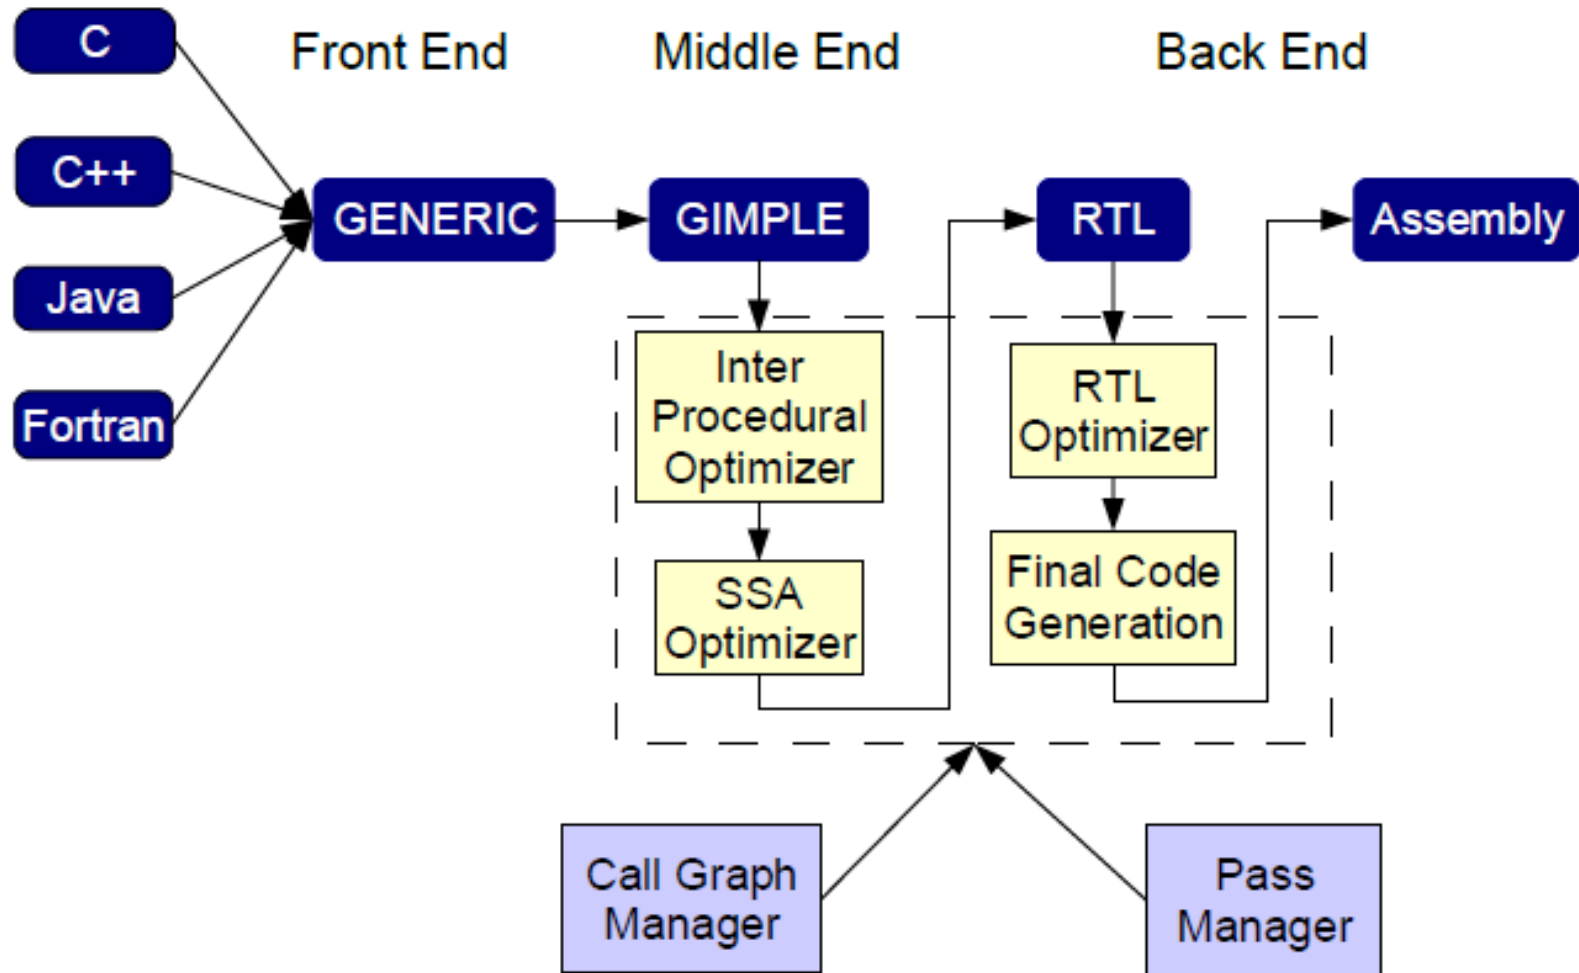
Structure of a Typical Compiler



GCC



GCC Compiler Pipeline



GCC – SSA Optimizers

- Operates on GIMPLE
- Around 100 Compiler Passes
 - Vectorization
 - Various loop optimizations
 - Traditional Scalar Optimizations: CCP, DCE, DSE, FRE, PRE, VRP, SRA, jump threading, forward propagation
 - Field-sensitive Points-to alias analysis
 - Pointer checking instrumentation for C/C++
 - Interprocedural analysis and optimizations: CCP, inlining, points-to analysis, pure/const and type escape analysis

GCC: RTL Optimizations

- Around 70 passes
- Operates closer to the target
 - Register Allocation
 - Scheduling
 - Software Pipelining
 - Common Subexpression Elimination
 - Instruction Recombination
 - Mode Switching Reduction
 - Peephole optimizations
 - Machine specific reorganization

GCC: GENERIC and GIMPLE IRs

GENERIC

```
if (foo (a + b, c))  
    c = b++ / a  
endif  
return c
```

High GIMPLE

```
t1 = a + b  
t2 = foo (t1, c)  
if (t2 != 0)  
    t3 = b  
    b = b + 1  
    c = t3 / a  
endif  
return c
```

Low GIMPLE

```
t1 = a + b  
t2 = foo (t1, c)  
if (t2 != 0) <L1, L2>  
L1:  
    t3 = b  
    b = b + 1  
    c = t3 / a  
    goto L3  
L2:  
L3:  
    return c
```

Open 64 Compiler Structure

