

# Detectors and Descriptors

P J Narayanan

**CS5765. Computer Vision. Spring 2013**

CVIT, IIT, Hyderabad



# Recognition Strategies


- Problem: Distinguish between day and night images.  
Or, distinguish between outdoor and indoor images
  - Solution: Use global features describing the image.  
Such as colour histogram, amount of edges, etc.
  - *Is this document about cricket?*
- Problem: Recognize special objects like faces, person.
  - Solution: Design specialized detectors
  - *Mark sentences about a wicket falling*
- We want to search **our keywords** on Google
- Can we look for *local* information (like words in documents) for search and/or recognition?

# Flexible Recognition/Retrieval

- How do we do more flexible recognition?
- Applications: Several
  - Searching for similar images. Give an image as the query to a search engine!
  - Distinguish between scenarios. Say chair images from tables. Or between individuals, etc.
- Strategy: find **interest points** and **descriptors**
  - Find “interesting” and “unique” features in image
  - Use their combination for high level reasoning and recognition.
- Like words in a document, enabling key-word search


# Interest Points: Requirements

- Repeatable across different transformations, noise, color change, etc.
- Informative or discriminative towards the task. Should capture the essential information relevant to the task
- Efficient in terms of computation time and space




# Interest Points

- What are the **important** elements in an image?
- Extract important elements and reason using them



# Uniform Regions, Edges, and Corners

- Which points are **distinct** in the *local neighbourhood*?
- Which points are **distinct** in the *whole image*?





# Detector & Descriptor

- **Interest Point Detector:** Where is the interest point?
  - Search the image and identify locations that qualify
  - What to look for? Big change in appearance that can be identified reliably across imaging scenarios
- **Descriptor:** What are the properties?
  - How do we describe it in a repeatable manner?
  - What properties to use? Color and appearance? Gradients or a measure of local changes? Texture pattern?

# Corners

- Regions where movement in all direction cause serious change in appearance
- How does the appearance change with a small displacement  $\vec{u} = [\Delta x \ \Delta y]^T$  at  $(x, y)$ ?
- $$D_{\vec{u}}(x, y) = \sum_{x, y \in W} w(x, y) [I(x + \Delta x, y + \Delta y) - I(x, y)]^2$$

where  $w(x, y)$  is a weighting function to give more importance to the centre. Gaussian in 2D works fine.
- Using Taylor expansion:
$$I(x + \Delta x, y + \Delta y) = I(x, y) + I_x(x, y)\Delta x + I_y(x, y)\Delta y + \dots$$



- Ignoring the quadratic terms and substituting in  $D(\cdot)$

$$\begin{aligned} D_{\vec{u}}(x, y) &= \sum_{x, y} \sum_{W} w(x, y) [I_x(x, y)\Delta x + I_y(x, y)\Delta y]^2 \\ &= \sum_{x, y} \sum_{W} w(x, y) [\Delta x \quad \Delta y] \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} \end{aligned}$$

- $D_{\vec{u}}(x, y) = [\Delta x \quad \Delta y] \mathbf{A} [\Delta x \quad \Delta y]^T$  where

$$\mathbf{A} = \sum_{x, y} \sum_{W} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = \begin{bmatrix} \sum_W I_x^2 & \sum_W I_x I_y \\ \sum_W I_x I_y & \sum_W I_y^2 \end{bmatrix}$$


- $D(\cdot)$  should be high in all directions  $\vec{u}$ .
- Examine the eigenvalues of  $\mathbf{A}$ . If both are high,  $D(\cdot)$  will be high in all directions

# Eigenanalysis

- $A = [e_1 \ e_2] \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} [e_1 \ e_2]^T = [e_1 \ e_2] \Lambda [e_1 \ e_2]^T$
- $x^T A x$  is an ellipse in 2D.
- Eigenvectors define an orthonormal (rotation) matrix.
- $x^T A x = ([e_1 \ e_2]^T x)^T \Lambda ([e_1 \ e_2]^T x) = (R x)^T \Lambda (R x)$
- Eigenvectors define a rotation matrix.  $R x$  orients the ellipse and the eigenvalues give the extents of its major and minor axes.

# Harris Corner

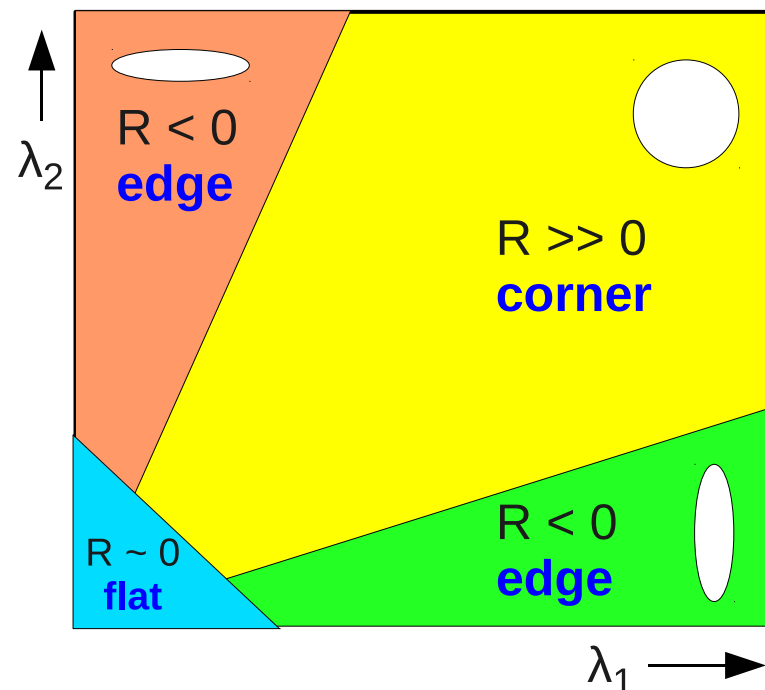
- Get eigenvalues of the Harris matrix  $A$  at each location.
- Regions with large variation in all directions: High  $\lambda_1$  and  $\lambda_2$
- Regions with large variation in one direction: High  $\lambda_1$  or  $\lambda_2$
- Near constant regions: Low  $\lambda_1$  and  $\lambda_2$
- We know that:  $\text{trace}(A) = \lambda_1 + \lambda_2$  and  $\det(A) = \lambda_1 \lambda_2$
- Define **cornerness**  $R = \det(A) - k \text{trace}(A)^2$ . Harris used  $k$  in the range  $[0.04, 0.06]$
- $R < 0$  when only one eigenvalue is strong.  $R \gg 0$  for corners, when both eigenvalues are strong



If  $\lambda_1 = r\lambda_2$ ,  $\frac{\det(A)}{\text{trace}(A)^2} = \frac{r\lambda_2^2}{(1+r)^2\lambda_2^2} = \frac{r}{(1+r)^2}$ , a function of  $r$  only.

# Corner Detection Algorithm


- Evaluate derivatives  $I_x, I_y$ , matrix  $A$  at all points based on a window and a weight function  $w(x, y)$
- Evaluate  $R$  everywhere.
- Corner if  $R > \tau$
- Find the maximum  $R$  in a  $3 \times 3$  or  $5 \times 5$  neighbourhood and declare it as a corner point.




$$A(\sigma_I, \sigma_D) = G(0, \sigma_I) \star \begin{bmatrix} I_x^2(\sigma_D) & I_x(\sigma_D)I_y(\sigma_D) \\ I_x(\sigma_D)I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

# Different Stages



Image:



Gradients:



Squares/Products:



Weighted Squares:



Corners:



# Corners on an Image



# Other Corner Detectors




- Hessian
- Laplacian, Difference of Gaussian (DoG)
- Hessian-Affine, Harris-Affine
- Hessian-Laplace, Harris-Laplace
- Maximally Stable Extremal Regions (MSER)
- And others ...



# Corners at Different Scales

- Corners and surrounding regions are not always of same size/scale
- Need to look for corners in different scales
- Ideally, only one “natural” scale should succeed
- Laplacian of Gaussian (LoG): a blob detector
- Characteristic scale: Where response to LoG is maximum in scale space

# LoG and Blob Detection





- Laplacian:  $\nabla^2 f = f_{xx} + f_{yy}$
- Laplacian of Gaussian:  $LoG(x, y, \sigma) = \nabla^2 G(x, y, \sigma)$


$$\nabla^2 G(x, y, \sigma) = \frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} e^{-(x^2 + y^2)/2\sigma^2}$$

- LoG is a filter that can be applied everywhere
- LoG responds to dark/bright **blobs** matching its scale.
- Change scale by changing the  $\sigma$  of the Gaussian

# LoG at Different Scales



Scale →



Scale →

# Corners & Characteristic Scale



- Scale signature: Apply LoG at different scales at a point and plot the response against the scale.
- The response peaks at the **characteristic scale** of the feature
- The blob feature makes most sense at the characteristic scale.
- Corner detection (like Harris corners) should also recover the scale of the feature.
- Find the local maximum in an  $x$ - $y$ - $\sigma$  space.

# SIFT [Lowe 1999, 2001]

- **Scale Invariant Feature Transform**: An influential detector and descriptor by David Lowe.
- Detects the scale, orientation of the feature and gives a gradient-based descriptor for it.
- Translation, rotation, and scale invariant interest points.
- Useful for image matching, recognition, robotics, etc.
- Four steps:
  1. Find the extrema in the scale space
  2. Detect the keypoint location
  3. Assign the dominant orientation
  4. Generate the gradient-based descriptor

# Difference of Gaussian (DoG)

- LoG gives the best results for the characteristic scale, but is expensive to compute.
- Lowe showed the *Difference of Gaussians (DoG)* does well. It approximates the scale-normalized LoG.
- Subtract Gaussians in different scales to get DoG




# Gaussian and Laplacian Pyramid

- DoG: Build a Gaussian Pyramid by convolving with Gaussians:  $L(u, v) = G(x, y, k\sigma) \star I(u, v)$ , for different values of  $k$
- Subtract to get  
 $D(u, v, k) = (G(x, y, k\sigma) - G(x, y, \sigma)) \star I(u, v)$
- Laplacian Pyramid: Lowest resolution Gaussian pyramid plus the differences at other resolutions.
- Typical value of  $k = 1.2 - 1.3$ . Make  $s$  levels in an octave where  $s = \log_k 2$  or  $k = 2^{1/s}$ .
- For the next octave, downsample the image by a factor of 2 and repeat.
- David Lowe found 3 scales per octave ( $k = 2^{\frac{1}{3}}$ ) and  $\sigma = 1.6$  to work best based on experiments.




# Scale Space for SIFT



# Scale Space Extrema Detection

- Consider the 3D neighbourhood of each pixel with 8 pixels in the current scale and 9 each in scales above and below.
- Declare the pixel a candidate interest point if its response is larger or smaller than all neighbours.
- Remove unstable ones: Those of *low contrast* and those *near edges*




# Orientation Assignment

- Given the selected (nearest) scale of each candidate key point, find the gradient magnitude and direction at the pixel in the Gaussian image of that scale.
- Form a 36-bin orientation histogram of pixels around the candidate point. Each pixel adds its Gaussian-weighted gradient magnitude to its orientation bin.
- Peak in the histogram represents the dominant orientation. This orientation is assigned to the keypoint.
- If another bin has a peak at 80% of the maximum, an additional keypoint with the new orientation is also generated.

# SIFT Descriptor


- Find gradient magnitude & direction at pixels around keypoint. Apply Gaussian weights to gradient magnitude.
- Rotate the directions with respect to the dominant direction to achieve rotation invariance.
- Find an orientation histogram of 8 bins and accumulate the weighted magnitudes in each.
- $16 \times 16$  regions, divided into 16 blocks of  $4 \times 4$ . Combine histograms into a vector. **Normalize.**

**SIFT descriptor:** A 128-dimensional ( $4 \times 4 \times 8$ ) vector of gradient orientation histograms






# An Example



# SIFT for Recognition



- Find SIFT keypoints and detectors in the image
- Can get sufficient matches for recognition even with a lot of occlusions.
- Use *ratio test* for matching a feature from one image to features in the other!



# SIFT for Panorama





# SIFT for Panorama



Matching: Find closest 2 neighbours in SIFT space.  
If ratio of their distances is low (say  $< 0.7$ ), correct match!!  
RANSAC to compute homography


# SIFT and Others

- SIFT is reasonably invariant to translation, rotation, scaling, intensity variations, and some affine transformations.
- Efficient implementations from David Lowe, VLFeat, etc.
- PCA-SIFT: Reduce dimensionality using PCA. It is not clear this benefits, however.
- SIFT is somewhat slow to compute, however.
- SIFT is also patented!!

# Other Descriptors

- SURF: Speeded-Up Robust Features.
  - Fast computation and detection.
  - Uses integral images.
  - Patent free!!
- Shape Context: Describes shapes by counting points in bins in a log-polar space
- Geometric Blur: Some photometric information along with shape
- Binary features: Local Binary Pattern (LBP), BRISK, BRIEF, ....
- GIST: A global feature, good for scene classification
- And many others....

# Local Binary Pattern (LBP)



- Compare the values of the middle pixel with its neighbours
- Encode greater as a “1” and lesser as a “0”
- An 8-bit binary number with 8 neighbours
- Can do at larger radius values also for different scales
- **The magnitude of the difference is no longer important. Only the sign is!**

# Detector vs Descriptor

- David Lowe proposed SIFT Detector and Descriptor
- These are not necessarily related. Detector finds good points to focus on. Descriptor says what goes on around them
- Detecting a keypoint followed by describing it is the usual order
- However, of late, the descriptor is evaluated at regular intervals in Computer Vision. Called **dense SIFT**
- Brings in sampling issues, etc., but is popular today

# Thank You!

P J Narayanan

Thanks to the Computer Vision community worldwide  
for many of the figures!!